

# Part VIII

## Transactions, Integrity and Triggers

# Transactions, Integrity and Triggers

## 1 Basic Terms

# Transactions, Integrity and Triggers

- 1 Basic Terms
- 2 Term Transaction

# Transactions, Integrity and Triggers

- 1 Basic Terms
- 2 Term Transaction
- 3 Possible Problems with Parallel Transactions

# Transactions, Integrity and Triggers

- 1 Basic Terms
- 2 Term Transaction
- 3 Possible Problems with Parallel Transactions
- 4 Transactions in SQL

# Transactions, Integrity and Triggers

- 1 Basic Terms
- 2 Term Transaction
- 3 Possible Problems with Parallel Transactions
- 4 Transactions in SQL
- 5 Integrity Constraints in SQL

# Transactions, Integrity and Triggers

- 1 Basic Terms
- 2 Term Transaction
- 3 Possible Problems with Parallel Transactions
- 4 Transactions in SQL
- 5 Integrity Constraints in SQL
- 6 Trigger

# Transactions, Integrity and Triggers

- 1 Basic Terms
- 2 Term Transaction
- 3 Possible Problems with Parallel Transactions
- 4 Transactions in SQL
- 5 Integrity Constraints in SQL
- 6 Trigger
- 7 Integrity Enforcement with Triggers



# Transactions, Integrity and Triggers

- 1 Basic Terms
- 2 Term Transaction
- 3 Possible Problems with Parallel Transactions
- 4 Transactions in SQL
- 5 Integrity Constraints in SQL
- 6 Trigger
- 7 Integrity Enforcement with Triggers
- 8 Trigger in Oracle

# Transactions, Integrity and Triggers

- 1 Basic Terms
- 2 Term Transaction
- 3 Possible Problems with Parallel Transactions
- 4 Transactions in SQL
- 5 Integrity Constraints in SQL
- 6 Trigger
- 7 Integrity Enforcement with Triggers
- 8 Trigger in Oracle
- 9 Summary

# Learning goals for today ...

- Understanding of fundamentals of integrity control in databases



# Learning goals for today ...

- Understanding of fundamentals of integrity control in databases
- Knowledge to formalize and implement integrity constraints



## Learning goals for today ...

- Understanding of fundamentals of integrity control in databases
- Knowledge to formalize and implement integrity constraints
- Knowledge of the transaction concept in databases



# Basic Terms

# Integrity

- **Integrity constraint** (*also: assertion*): Condition for the "permissibility" or "correctness"
- with respect to databases:
  - ▶ (single) database states,
  - ▶ state transitions from an old to a new database state,
  - ▶ long term database evolution

# Classification of Integrity

Constraint Class		Temporal Context
static		database state
dynamic	transitional temporal	state transition state sequence



# Inherent Integrity Constraints in the RM

- ① *Type Integrity:*
  - ▶ SQL allows domain definitions for a range of values for attributes
  - ▶ Permission or forbidding of null values
- ② *Key Integrity:*
  - ▶ Specification of a key for a relation
- ③ *Referential Integrity:*
  - ▶ Specification of foreign keys

# Term Transaction

# Example Scenarios

- Seat reservation for flights simultaneously from multiple travel agencies  
→ Seat could be sold multiple times when multiple travel agencies identify the seat as available
- Overlapping account operations of a bank
- Statistics database operations  
→ results are corrupted when data is changed during the calculation

# Transaction

A **transaction** is a sequence of operations (actions) that transforms the database from a consistent state into a consistent, possibly changed, state, while the **ACID-principle** must be hold.

- Aspects:

- ▶ Semantic Integrity: Correct (consistent) DB-state after a transaction has finished
- ▶ Operational Integrity: Prevent fault caused by "simultaneous" access of multiple users on the same data

# ACID-Properties

- **Atomicity:**

Transaction is executed completely or not at all

- **Consistency:**

Database is before the start and after the end of a transaction in a consistent state

- **Isolation:**

User, who is working on a database, should have the impression that she works alone on the database

- **Durability (Persistence):**

The result of transaction has to be saved "permanently" in a database after the transaction completed successfully

# Commands of a Transaction Language

- Begin of a transaction: Begin-of-Transaction-Command **BOT** (implicit in SQL!)
- **commit**: the transaction should try to finish successfully
  - ▶ success is not guaranteed!
- **abort**: the transaction has to be aborted
  - ▶ abort is guaranteed!

# Transaction: Integrity Violation

- Example:


- ▶ Transfer of an amount A from a household post K1 to another post K2
- ▶ Condition: Sum of the account balances stays constant

- Simplified notation

*Transfer* =  $\langle K1 := K1 - A; K2 := K2 + A \rangle;$

- Realization in SQL: as sequence of two elementary changes  $\rightsquigarrow$   
Condition is not necessarily fulfilled between single changing steps!

# Transaction: Behavior at System Crash





# Transaction: Behavior at System Crash /2

- **Consequences:**

- ▶ Contents of the volatile memory at the time  $t_f$  is unusable → transactions in different ways affected by this

- **Transaction states:**

- ▶ Still active transactions at the time of the failure ( $T_2$  and  $T_4$ )
- ▶ Already finished transactions before the time of the failure ( $T_1$ ,  $T_3$  and  $T_5$ )

# Simplified Model for Transactions

- Representation of database changes of a transaction
  - ▶ **read**( $A, x$ ): assign the value of the DB-object  $A$  to the variable  $x$
  - ▶ **write**( $x, A$ ): save the value of the variable  $x$  in the DB-object  $A$
- Example of a transaction  $T$ :

```
read( $A, x$ );  $x := x - 200$ ; write( $x, A$ );  
read( $B, y$ );  $y := y + 100$ ; write( $y, B$ );
```

- Execution variants for two transactions  $T_1, T_2$ :
  - ▶ serially, e.g.  $T_1$  before  $T_2$
  - ▶ "mixed", e.g. alternating steps of  $T_1$  and  $T_2$

# Possible Problems with Parallel Transactions

# Problems with Multi-User Operation

- Nonrepeatable Read
- Dependencies on not released data: Dirty Read
- The Phantom-Problem
- Lost Update

# Nonrepeatable Read

Example:

- Assurance  $x = A + B + C$  at the end of transaction  $T_1$
- $x, y, z$  are local variables
- $T_i$  is the transaction  $i$
- Integrity constraint  $A + B + C = 0$

# Example for Nonrepeatable Read

$T_1$	$T_2$
<b>read</b> ( $A, x$ );          <b>read</b> ( $B, y$ ); $x := x + y$ ; <b>read</b> ( $C, z$ ); $x := x + z$ ; <b>commit</b> ;	<b>read</b> ( $A, y$ ); $y := y/2$ ; <b>write</b> ( $y, A$ ); <b>read</b> ( $C, z$ ); $z := z + y$ ; <b>write</b> ( $z, C$ ); <b>commit</b> ;

# Dirty Read

$T_1$	$T_2$
<b>read</b> ( $A, x$ ); $x := x + 100$ ; <b>write</b> ( $x, A$ );  <b>abort</b> ;	<b>read</b> ( $A, x$ ); <b>read</b> ( $B, y$ ); $y := y + x$ ; <b>write</b> ( $y, B$ ); <b>commit</b> ;

# The Phantom-Problem

$T_1$	$T_2$
<pre><b>select count (*)</b> <b>into</b> X <b>from</b> Customer;  <b>update</b> Customer <b>set</b> Bonus = Bonus + 10000/X; <b>commit</b>;</pre>	<pre><b>insert</b> <b>into</b> Customer <b>values</b> ('Meier', 0, ...); <b>commit</b>;</pre>



# Lost Update

$T_1$	$T_2$	$A$
<b>read</b> ( $A, x$ );		10
	<b>read</b> ( $A, x$ );	10
$x := x + 1$ ;		10
	$x := x + 1$ ;	10
<b>write</b> ( $x, A$ );		11
	<b>write</b> ( $x, A$ );	11

# Serializability

An interleaved execution of multiple transactions is called **serializable**, if its effect is identical to the effect of a (arbitrarily chosen) serial execution of these transactions.

- Problem for checking serializability:
  - ▶ there are  $n!$  different serial execution orders for  $n$  transactions...
- **Schedule:** Plan of execution for transactions (ordered list of transaction operations)

# Transactions in SQL

# Transactions in SQL-DBS

## Weakening of ACID in SQL: Isolation levels

```
set transaction
  [ { read only | read write }, ]
  [isolation level
    { read uncommitted |
      read committed |
      repeatable read |
      serializable }, ]
  [ diagnostics size ...]
```

Default settings:

```
set transaction read write,
  isolation level serializable
```

# Meaning of Isolation Levels

## ● **read uncommitted**

- ▶ weakest level: access to not committed data, only for **read only** transactions
- ▶ statistic and similar transactions (approximate overview, incorrect values possible)
- ▶ no locks → efficient executable, other transactions are not hindered

## ● **read committed**

- ▶ only read finally written values, but *nonrepeatable read* possible

## ● **repeatable read**

- ▶ no *nonrepeatable read*, but phantom-problem can occur

## ● **serializable**

- ▶ guarantees serializability

# Isolation Levels: read committed

	$T_1$	$T_2$
	<b>set transaction isolation level read committed</b>	
1	<b>select</b> Name <b>from</b> WINES <b>where</b> WineID = 1014 → <i>Riesling</i>	
2		<b>update</b> WINES <b>set</b> Name = 'Riesling Superi- ore' <b>where</b> WineID = 1014
3	<b>select</b> Name <b>from</b> WINES <b>where</b> WineID = 1014 → <i>Riesling</i>	
4		<b>commit</b>
5	<b>select</b> Name <b>from</b> WINES <b>where</b> WineID = 1014 → <i>Riesling Superiore</i>	

# Isolation Levels: read committed /2

	$T_1$	$T_2$
	<b>set transaction isolation level read committed</b>	
1	<b>select</b> Name <b>from</b> WINES <b>where</b> WineID = 1014	
2		<b>update</b> WINES <b>set</b> Name = 'Riesling Superiore' <b>where</b> WineID = 1014
3	<b>update</b> WINES <b>set</b> Name = 'Superiore Riesling' <b>where</b> WineID = 1014 → <b>blocked</b>	
4		<b>commit</b>
5	<b>commit</b>	

# Isolation Levels: **serializable**

	$T_1$	$T_2$
	<b>set transaction isolation level serializable</b>	
1	<b>select</b> Name <b>into</b> N <b>from</b> WINES <b>where</b> WineID = 1014 → N := <i>Riesling</i>	
2		<b>update</b> WINES <b>set</b> Name = 'Riesling Superi- ore' <b>where</b> WineID = 1014
4		<b>commit</b>
5	<b>update</b> WINES <b>set</b> Name = 'Superior'    N <b>where</b> WineID = 1014 → <span style="border: 1px solid black; padding: 2px;">Abort</span>	



# Integrity Constraints in SQL

# Integrity Constraints in SQL-DDL

- **not null**: Null values prohibited
- **default**: Specification of default values
- **check ( *search-condition* )**: Attribute specific constraint (usually *One-Tuple-Integrity-Condition*)
- **primary key**: Specification of a primary key
- **foreign key ( *Attribute(e)* )**  
**references *Table( Attribute(e) )***:  
Specification of the referential integrity

# Integrity Constraints: Range of Values

- **create domain**: Establishing of a user defined range of values
- Example

```
create domain WineColor varchar(5)  
    default 'Red'  
    check (value in ('Red', 'White', 'Rose'))
```

- Application

```
create table WINES (  
    WineID int primary key,  
    Name varchar(20) not null,  
    Color WineColor,  
    ...)
```

# Integrity Constraints: **check-Clause**

- **check**: Establishing of further local integrity constraints within the defined range of values, attributes and relational scheme
- Example: Restriction of permitted values
- Example

```
create table WINES (  
    WineID int primary key,  
    Name varchar(20) not null,  
    Year int check(Year between 1980 and 2010),  
    ...  
)
```

# Preservation of Referential Integrity

- Checking of foreign keys after database changes
- for  $\pi_A(r_1) \subseteq \pi_K(r_2)$ ,  
e.g.  $\pi_{\text{Vineyard}}(\text{WINES}) \subseteq \pi_{\text{Vineyard}}(\text{PRODUCER})$ 
  - ▶ Tuple  $t$  is inserted into  $r_1 \Rightarrow$  check, whether  $t' \in r_2$  exists with:  
 $t'(K) = t(A)$ , d.h.  $t(A) \in \pi_K(r_2)$   
if not  $\Rightarrow$  reject
  - ▶ Tuple  $t'$  is removed from  $r_2 \Rightarrow$  check, whether  $\sigma_{A=t'(K)}(r_1) = \{\}$ , i.e. no tuple from  $r_1$  references  $t'$   
if not empty  $\Rightarrow$  reject or remove tuple from  $r_1$ , that reference  $t'$  (at cascading deletion)

# Checking Modes of Constraints

- **on update | delete**

Specification of a triggering event that starts the checking of the condition

- **cascade | set null | set default | no action**

**Cascading:** Handling of some integrity violations propagates over multiple levels, e.g. deletion as reaction on a violation of the referential integrity

- **deferred | immediate** sets the checking time for a condition

- ▶ **deferred:** put back to the end of the transaction
- ▶ **immediate:** immediate verification at any relevant database change

# Checking Modes: Example

- Cascading deletion

```
create table WINES (  
  WineID int primary key,  
  Name varchar(50) not null,  
  Price float not null,  
  Jahr int not null,  
  Vineyard varchar(30),  
  foreign key (Vineyard) references PRODUCER (Vineyard)  
    on delete cascade)
```

# The **assertion**-Clause

- Assertion: Predicate expressed by a condition that always has to be fulfilled by a database
- Syntax (SQL:2003)

```
create assertion name check ( predicate )
```

- Example:

```
create assertion Prices check  
  ( ( select sum (Price)  
      from WINES ) < 10000 )
```

```
create assertion Prices2 check  
  ( not exists (  
      select * from WINES where Price > 200) )
```



# Trigger

# Trigger

- Trigger: Statement/Procedure that is executed automatically by the DBMS at the occurrence of a specific event
- Application:
  - ▶ Enforcement of integrity conditions ("implementation" of integrity rules)
  - ▶ Auditing of DB-actions
  - ▶ Propagation of DB-changes
- Definition:

```
create trigger ...  
after <Operation>  
<Procedure>
```

# Example for Triggers

- Realization of a calculated attribute with two triggers:
  - ▶ Introduction of new tasks

```
create trigger TaskCounter+  
  on insertion of Task A:  
  update Customer  
  set NrTasks = NrTasks + 1  
  where CName = new A.CName
```

- ▶ Analogously for deletion of tasks:

```
create trigger TaskCounter-  
  on deletion ...:  
  update ...- 1 ...
```

# Trigger: Design and Implementation

- Specification of
  - ▶ *Event* and *condition* for activation of the trigger
  - ▶ *Action(s)* for the execution
- Syntax in SQL:2003 defined
- Available in most commercial systems (but with different syntax)

# SQL:2003-Trigger

- Syntax:

```
create trigger <Name:>  
after | before <Event>  
on <Relation>  
[ when <Condition> ]  
begin atomic < SQL-statements > end
```

- Event:

- ▶ **insert**
- ▶ **update** [ **of** <list of attributes> ]
- ▶ **delete**

## Further Specifications for Triggers

- **for each row** resp. **for each statement**: Activation of the trigger for *each* single change of a set-valued change or just once for the whole change
- **before** resp. **after**: Activation *before* or *after* the change
- **referencing new as** resp. **referencing old as**: Binding of a tuple variable on the new introduced resp. just removed ("old") tuple of a relation  
     $\rightsquigarrow$  tuple of the *difference relation*

## Example for Triggers

- *No customer account can fall below 0:*

```
create trigger bad_account
after update of Acc on CUSTOMER
referencing new as INSERTED
when (exists
    (select * from INSERTED where Acc < 0)
)
begin atomic
    rollback;
end
```

↪ similar trigger for **insert**

## Example for triggers /2

- Producers **must** be removed, if they do not offer any wine:

```
create trigger useless_Vineyard
after delete on WINES
referencing old as o
for each row
when (not exists
    (select * from WINES W
     where W.Vineyard = o.Vineyard))
begin atomic
    delete from PRODUCER where Vineyard = o.Vineyard;
end
```



# Integrity Enforcement with Triggers

# Integrity Enforcement with Triggers

1. Specify object  $o_i$ , for which the condition  $\phi$  should be monitored
  - ▶ Usually monitor multiple  $o_i$  when condition is across relations
  - ▶ Candidates for  $o_i$  are tuples of the relation names that occur in  $\phi$
2. Specify the elemental database changes  $u_{ij}$  on objects  $o_i$  that can violate  $\phi$ 
  - ▶ Rules: e.g., check existence requirements on deletion and updates, but not on insertion etc.

## Integrity Enforcement with Triggers /2

3. Specify, depending on the application, the reaction  $r_i$  on the integrity violation
  - ▶ Reset the transaction (**rollback**)
  - ▶ Correcting database changes
4. Formulate following triggers:

```
create trigger t-phi-ij after  $u_{ij}$  on  $o_i$   
when  $\neg\phi$   
begin  $r_i$  end
```

5. If possible, simplify the created trigger

# Trigger in Oracle

# Trigger in Oracle

- Implementation in PL/SQL
- Notation

```
create [ or replace ] trigger trigger-name  
before | after  
insert or update [ of columns ]  
or delete on table  
[ for each row  
[ when ( predicate ) ] ]  
PL/SQL-Block
```

# Trigger in Oracle: Types

- Statement level trigger: Trigger is triggered before resp. after the DML-statement
- Row level trigger: Trigger is triggered before resp. after each single modification (*one tuple at a time*)

Trigger on row level:

- Predicate for restriction (**when**)
- Access on old (**:old.col**) resp. new (**:new.col**) tuple
  - ▶ for **delete**: only (**:old.col**)
  - ▶ for **insert**: only (**:new.col**)
  - ▶ in **when**-clause only (**new.col**) resp. (**old.col**)

## Trigger in Oracle /2

- Transaction abortion with **raise\_application\_error**(*code*, *message*)
- Distinction of the type of the DML-statement

```
if deleting then ... end if;  
if updating then ... end if;  
if inserting then ... end if;
```

# Trigger in Oracle: Example

- *No customer account can fall below 0:*

```
create or replace trigger bad_account
after insert or update of Acc on Customer
for each row
when (:new.Acc < 0)
begin
    raise_application_error(-20221,
        'Not below 0');
end;
```



# Summary

# Summary

- Enforcement of correctness resp. integrity of the data
- Inherent integrity constraints of the relational model
- Additional SQL-integrity constraints: **check**-clause, **assertion**-statement
- Trigger for "implementation" of integrity constraints resp. rules

# Control Questions

- What is the purpose of integrity enforcement? Which types of integrity constraints are there?



# Control Questions

- What is the purpose of integrity enforcement? Which types of integrity constraints are there?
- How can integrity constraints and rules be formulated in SQL systems?



# Control Questions

- What is the purpose of integrity enforcement? Which types of integrity constraints are there?
- How can integrity constraints and rules be formulated in SQL systems?
- What requirements result from the ACID-principle? How are these achieved in database systems?

